



Performance Programming with Fortran

Jeremy Roberts

01/31/2012

Assessing Performance

Increasing Performance: Leveraging Libraries

Increasing Performance: Multithreading

Summary

1. C. Evangelinos, 12.950 course slides
2. Chapman, B., *et al.*, **Using OpenMP**, The MIT Press (2008)
3. www.openmp.org/mp-documents/OpenMP3.0-FortranCard.pdf
4. <https://computing.llnl.gov/tutorials/openMP/>

There is a ton of information out there for all the stuff in these slides!

Go get the Intel compilers (for noncommercial use) at:
<http://software.intel.com/en-us/articles/non-commercial-software-download/>

Bare code in `/home/fortran12/users/robertsj`

What defines a program's performance?

Possible metrics:

1. wall time
2. CPU time
3. FLOPS (floating-point **o**perations **p**er **s**econd)
4. memory efficiency (*i.e.* minimizing cache hits)
5. parallel scaling
6. (maintainability)
7. (correctness)

Others?

An Example: Matrix-Vector Multiplication

We want to compute

$$\vec{y} \leftarrow \mathbf{A}\vec{x},$$

where \mathbf{A} is an $n \times n$ matrix, and x and y are n -vectors. *Conceptually*, this is a bunch of dot products

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} \vec{a}_{1:} \cdot \vec{x} \\ \vec{a}_{2:} \cdot \vec{x} \\ \vec{a}_{3:} \cdot \vec{x} \end{bmatrix}$$

or (better) a linear combination of the columns of \mathbf{A}

$$\mathbf{A}\vec{x} = x_1\vec{a}_{:1} + x_2\vec{a}_{:2} + x_3\vec{a}_{:3}.$$

Numerically, we compute

$$y_i = \sum_{j=1}^n a_{ij}x_j, \quad i = 1, \dots, n.$$

Listing 1: matvec-driver-bare.f90

```
1  ! Driver program to test matrix-vector multiplication
2  program matvec_driver
3      implicit none
4      double precision, allocatable :: A(:, :)
5      double precision, allocatable :: x(:), y(:)
6      double precision :: alpha = 1.0, beta = 0.0
7      integer :: n, m, lda, incx = 1, incy = 1, i, num_loops
8      character*1 :: trans = 'n'
9      m = 2000
10     n = m
11     lda = m
12     ! Create the matrix and vector
13
14     ! Initialize A and x
15
16     ! Loop over and apply A several times for consistent timing
17     num_loops = 100
18     do i = 1, num_loops
19         ! Reset
20         y = 0.0
21         call dgemv('n', m, n, alpha, A, lda, x, incx, beta, y, incy)
22     end do
23
24 end program matvec_driver
```

Fill it out and name it matvec-driver.f90.

Listing 2: matvec-bare.f90

```
1  subroutine dgemv(trans, m, n, alpha, A, lda, x, incx, beta, y, incy)
2  ! Simplified DGEMV that does
3  !   y := A*x
4  ! We follow the BLAS signature for compatibility.  For more,
5  ! see http://netlib.org/blas/.
6  implicit none
7  ! input/output
8  character*1, intent(in)      :: trans
9  integer, intent(in)          :: m, n, lda, incx, incy
10 double precision, intent(in)  :: alpha, beta
11 double precision, intent(in)  :: A(m, n)
12 double precision, intent(in)  :: x(n)
13 double precision, intent(inout) :: y(m)
14
15
16
17
18
19
20
21
22 end subroutine dgemv
```

Implement $y_i = \sum_{j=1}^n a_{ij}x_j$, $i = 1, \dots, n$, and name it matvec-first.f90.

The Simplest Timer: `time`

First, compile

```
gfortran -o matvec-first \  
matvec-first.f90 matvec-driver.f90
```

Then on a Linux machine (maybe Macs, too), run

```
time matvec-first
```

and get

```
real      0m4.511s  
user      0m4.470s  
sys       0m0.020s
```

The user + sys time is **CPU time**. The real time is **wall time**.

Question: $t_{CPU} \geq t_{wall}$ or $t_{CPU} \leq t_{wall}$?

Listing 3: wtime.f90

```
1 double precision function wtime()
2 ! WTIME returns a reading of the wall clock time. Use as follows
3 !
4 !   double precision :: t
5 !   t = wtime()
6 !   ! do some work
7 !   print *, "elapsed time = ", wtime() - t, " seconds."
8 !
9 ! This is a slight modification of John Burkardt's function of the same
10 ! name. See http://people.sc.fsu.edu/~jburkardt for this and other useful
11 ! source.
12 implicit none
13 integer count          ! processor-dependent value based on current value
14                        ! of processor clock
15 integer count_rate     ! clock counts per second
16 integer count_max     ! maximum value that count can take
17 call system_clock(count, count_rate, count_max)
18 wtime = dble(count) / dble(count_rate)
19 end function wtime
```

Add this timer to the matvec-driver.f90.

Compiler-based Timer

Compile

```
gfortran -o matvec-first wtime.f90 \
    matvec-first.f90 matvec-driver.f90
```

and then run

```
matvec-first
```

and get

```
elapsed time =      4.4969999999739230      seconds.
```

Which time is this? Also, note the poor precision. Apparently gfortran's clock isn't too precise.

Other timing utilities are provided by

1. OpenMP (`omp_get_wtime()`; see later slides)
2. MPI (`MPI_Wtime()`)
3. Hardware vendors
4. Software vendors

For most scientific applications, the OpenMP and MPI functions should suffice; YMMV.

How many floating point operations in \mathbf{Ax} ? (Assume square.)

$$y_i = \sum_{j=1}^n a_{ij}x_j, \quad i = 1, \dots, n.$$

```
do i = 1, n
  do j = 1, n
    y(i) = y(i) + A(i, j) * x(j)
  end do
end do
```

For $m = n = 2000$ and 100 runs, $t_{CPU} \approx 4.5$ seconds. How many FLOPS?

Add a line to compute FLOPS in `matvec-driver.f90`.

Linux users: learn about your machine. Execute `sudo lshw > out.txt`. Digging through, I find

```
*-cpu
  description: CPU
  product: Intel(R) Core(TM) i7 CPU           970  @ 3.20GHz
...
  description: L1 cache
...
  slot: L1-Cache
  size: 192KiB
...
  description: L2 cache
...
  slot: L2-Cache
  size: 1536KiB
...
  description: L3 cache
...
  slot: L3-Cache
  size: 12MiB
```

So I crunch $3.2 \cdot 10^9$ times per second. That means I should get 3.2 GFLOPS, right?

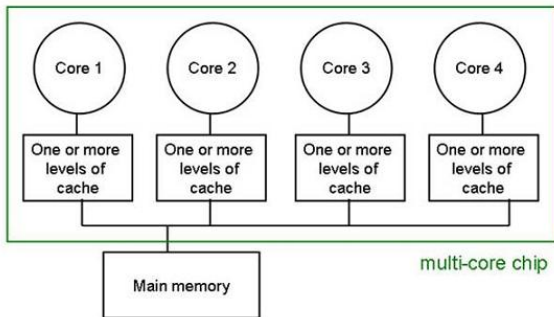
Optimize:

```
gfortran -O3 -o matvec-first matvec-first.f90 matvec-driver.f90
matvec-first
elapsed time =      3.5709999999962747      seconds.
```

or about 225 MFLOPS. Still over an order of magnitude off. Another compiler?

```
ifort -O3 -o matvec-first matvec-first.f90 matvec-driver.f90
matvec-first
elapsed time =      0.3679999999987660      seconds.
```

or about 2.2 GFLOPS. Wow! But not yet optimal.



Now remember L1, L2, and L3 cache sizes (192 KiB, 1536 KiB, and 12 MB). A double precision number is 8 bytes. For $n = 2000$, one vector is only 16 KiB but the matrix is 32 MB! **What's our issue?**

Fortran uses *column-major* storage. Thus, the elements of matrix

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$

are stored as $a_{11}, a_{21}, a_{31}, a_{12}$, and so on. But look at

```
do i = 1, n
  do j = 1, n
    y(i) = y(i) + A(i, j) * x(j)
  end do
end do
```

Implement our fix, naming it `matvec-second.f90`.

FLOPS - A Fix?

Switch order:

```
do j = 1, n      ! j and i switched
  do i = 1, n
    y(i) = y(i) + A(i, j) * x(j)
```

Compile and run

```
gfortran -O3 -o matvec-second matvec-second.f90 matvec-driver.f90
you@pc:~/fpp_code/$ matvec-second < input.txt
elapsed time = 0.38800000003539026 seconds.
```

which is much better for gfortran, but

```
ifort -O3 -o matvec-first matvec-second.f90 matvec-driver.f90
you@pc:~/fpp_code/$ matvec-second < input.txt
elapsed time = 0.374899999995250 seconds.
```

is the same as before! Why? Are we optimal yet?

Setting 1000 loops and using AMD's Math Core Library:


```
gfortran -O3 -o matvex-amcl wtime.f90 matvec-driver.f90
-L /home/robertsj/opt/acml/acml4.4.0/gfortran64/lib/
-lacml -lgfortran
matvec-amcl < input.txt
elapsed time =      3.4669999999459833          seconds.
```

again, about 2.25 GFlops. Using Intel's Math Kernel Library:

```
ifort -O3 -o matvex-mkl wtime.f90 matvec-driver.f90 -i8
-ISMKLROOT/include -Wl,--start-group
$MKLROOT/lib/intel64/libmkl_intel_ilp64.a
$MKLROOT/lib/intel64/libmkl_sequential.a
$MKLROOT/lib/intel64/libmkl_core.a -Wl,--end-group -lpthread -lm
matvec-mkl < input.txt
elapsed time =      3.05934691429138          seconds.
```

or about 2.61 GFlops.

Only Masochists Implement Linear Algebra Themselves! Leave it to the pros!



Problem and cache size are important: *the CPU can only do its fastest work if it gets the data just as fast.*

Suggested Homework: Adapt `matvec-driver.f99` to produce a graph like the one shown.


If you do numerical anything, you might run into these:

- ▶ BLAS, LAPACK, *etc.*: **standard** linear algebra libraries with many high performance vendor implementations (AMD, Intel, ATLAS)
- ▶ PETSc: parallel linear and nonlinear solvers
- ▶ SLEPc: parallel eigensolvers
- ▶ Trilinos: parallel solvers
- ▶ MOOSE: framework for multiphysics PDE's
- ▶ HDF5: library for scientific data

... plus many, many more. Don't reinvent the wheel!

The OpenMP Programming Model

- ▶ OpenMP is a standard defining an API for *shared memory* applications
- ▶ Employs the “fork-join” model
- ▶ Encourages incremental parallelization
- ▶ Fortran (original), C/C++



When performing work in parallel

- ▶ threads read and write *shared* data
- ▶ threads can read and write private copies
- ▶ synchronization is required to avoid “race conditions”
- ▶ loop-carried dependencies must be avoided

Note, memory is multi-level: cache is *not* shared. One must be careful to make sure all threads see what they should see.

Example: Computing π in Parallel

Listing 4: pi-bare.f90

```

1  ! Computing pi in parallel by numerical integration
2  program pi
3      use omp_lib
4      implicit none
5      double precision :: pie, total, x, dx, t
6      integer          :: i, number_steps
7
8  end program pi
  
```

Implement the integration and some print statements. Name it pi-first.f90

See the OpenMP crib sheet.

Example: Computing π in Parallel

```
gfortran -O3 -fopenmp pi-first.f90 -o pi-first
export OMP_NUM_THREADS=4
./pi-first
```

and get

```
hi from thread          0 of          4
hi from thread          1 of          4
hi from thread          2 of          4
hi from thread          3 of          4
      pi =      2.2703528624655167
elapsed time =      0.17734825599472970      seconds.
```


Whoops! Our first mistake in parallel programming. It will happen again.

Now we get

```
gfortran -O3 -fopenmp pi-first.f90 -o pi-first
export OMP_NUM_THREADS=1
./pi-second
hi from thread          0 of          1
    pi =    3.1415926453617899
elapsed time =    0.66723693604581058      seconds.
export OMP_NUM_THREADS=2
./pi-second
hi from thread          0 of          2
hi from thread          1 of          2
    pi =    3.1415926469906479
elapsed time =    0.34848214697558433      seconds.
export OMP_NUM_THREADS=4
./pi-second
hi from thread          0 of          4
hi from thread          3 of          4
hi from thread          2 of          4
hi from thread          1 of          4
    pi =    3.1415926454477359
elapsed time =    0.17908076802268624      seconds.
```

Look at that *weak scaling*! **Suggested homework: test its *strong scaling*.**

Add OpenMP to the matvec subroutine.



Curiously, the best approach is to switch back i and j ! See the text by Chapman et al. for details. (Note, things like MKL have threaded versions, too, so still avoid reinventing the wheel.)

The Monte Carlo Slab Problem

Homework (ha!): parallelize your program.

Given: all the code from yesterday plus a parallel-ready random number generator (if not already provided).

Solution Sketch: The solution looks a lot like the π problem above. You need to

- ▶ wrap the problem in a `parallel` directive
- ▶ allocate and initiate local tally arrays
- ▶ wrap the histories in a `do` directive
- ▶ note the global object `tal` is already given the attribute `threadprivate` (like `private` but for things in a potentially global scope such as modules)
- ▶ carefully select which variables are denoted `private` (e.g. the `mat` variable) and which are `public`
- ▶ *safely* add the local tally information to the global `tal` variable. Hint: look up the `critical` directive.

Note: I changed roughly 15 lines of code to parallelize the serial code, and most of this was in replacing `tal` with a `private` variable.

Key takeaways:

- ▶ Speed isn't the be all end all; memory is often limiting
- ▶ The compiler and library choice is important for peak performance
- ▶ Don't reinvent the wheel (in particular, OMILAT)
- ▶ OpenMP is easy (mostly)
- ▶ Computing is fun!

Help is available if you decide to attempt the problem...